

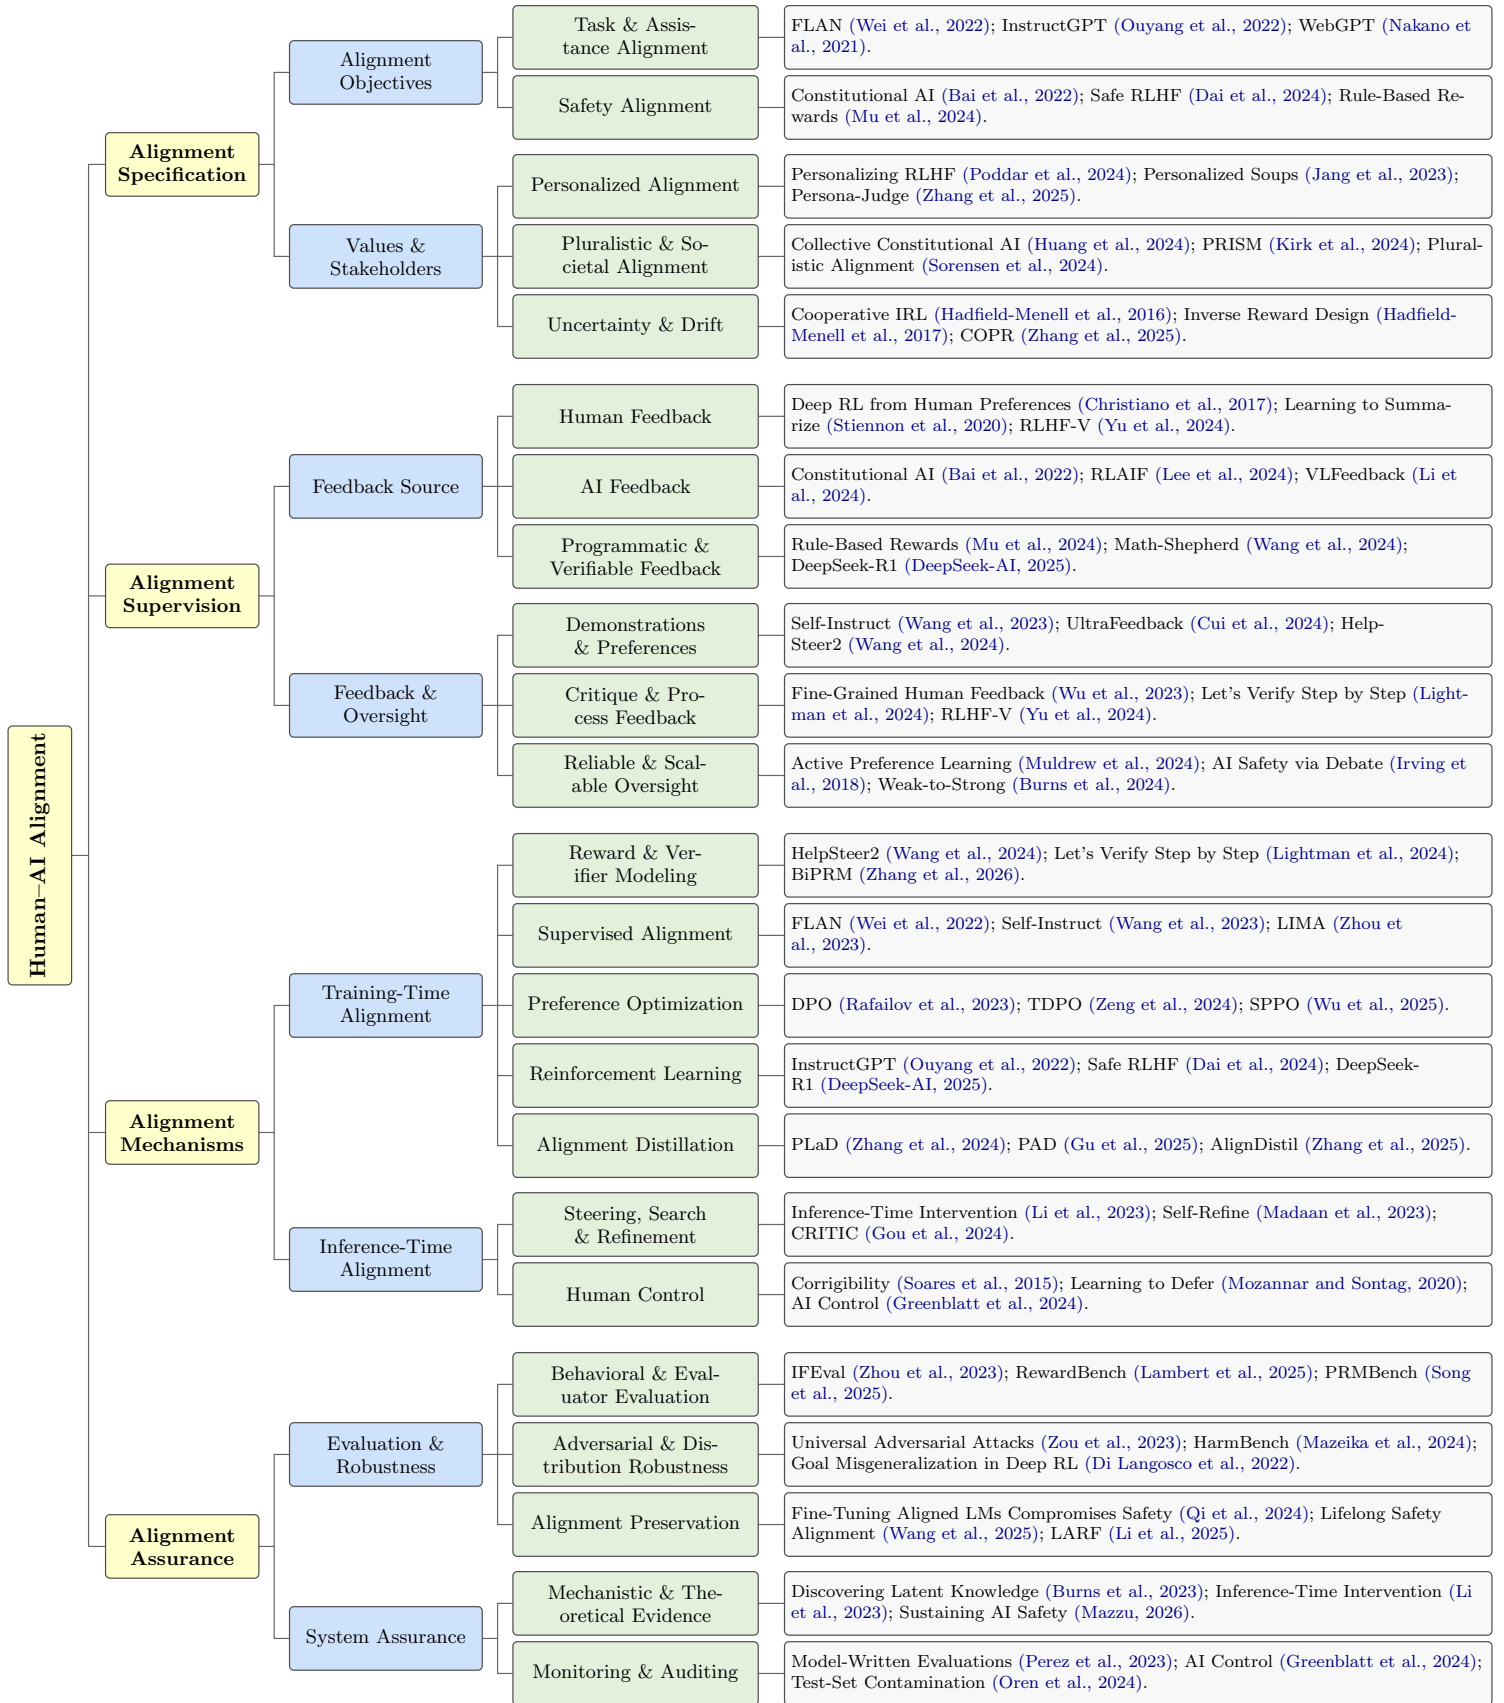


Figure 1: A compact taxonomy of Human-AI Alignment organized along four complementary dimensions and 23 terminal categories. Each gray box lists three illustrative references, while the accompanying mapping retains the full bibliography. Categories are non-exclusive. Demonstrations and preferences include pairwise comparisons, scalar ratings, and multi-attribute judgments. Mechanistic and theoretical evidence is interpreted under the assumptions of the corresponding analyses, while process-verifier quality is assessed within behavioral and evaluator evaluation.